



SDK White Paper

Occlusion Query Checking for Hidden Pixels


WP-01402-001_v01
July 2004

nVSDK

Occlusion Queries

Many graphics engines today waste time by attempting to draw pixels of objects that do not appear in a scene. These hidden objects are *occluded* from view, but often the engine rendering a scene does not know which pixels are occluded and which are visible.

Tracking information about which pixels are occluded in a scene and which are visible is not a trivial design problem. Because metadata about the relation of objects in a scene is not available, most modern hardware supports the concept of an “occlusion query.” Occlusions queries ask the hardware to report on whether or not a selection of the frames rendered were occluded when the buffer is presented to the screen.



Occlusion queries provide useful information that can be used in a variety of ways to reduce the load on the graphics card. Strategies for using the information include not drawing objects that are not visible and reducing other non-graphics related loads, such as scaling down physics and AI for objects that are not visible.

The demonstration described in this document presents a technique for not drawing an extremely complex object once it becomes completely occluded.

For more information, contact

Bryan Dudash
bdudash@nvidia.com

NVIDIA Corporation
2701 San Tomas Expressway
Santa Clara, CA 95050

D3D IDirect3DQuery9

The D3D API provides an interface, called IDirect3DQuery9, which gives access to an occlusion query. Each call to IDirect3DQuery9 represents a single outstanding query. This section briefly describes the IDirect3DQuery9 API. However, please consult the DirectX9 documentation for more information.

IDirect3DQuery9 API for Occlusion Queries

```
HRESULT IDirect3DDevice9::CreateQuery(D3DQUERYTYPE Type,  
IDirect3DQuery9** ppQuery);
```

To make an occlusion query, pass in D3DQUERYTYPE_OCCLUSION as the type, and a pointer to a query interface object. From then on, the query object interface is all you'll use.

```
HRESULT IDirect3DQuery9::Issue(DWORD dwIssueFlags);
```

Wrap the render calls that you are interested in monitoring with calls to Issue(D3DISSUE_BEGIN) and Issue(D3DISSUE_END).

```
HRESULT IDirect3DQuery9::GetData(void* pData, DWORD dwSize,  
DWORD dwGetDataFlags);
```

When calling GetData() on an occlusion query. The pData pointer should point to a DWORD. dwSize should be sizeof(DWORD). When called with no flags, the function will return immediately. It will return S_OK if the query results have returned. If not, then the query is still outstanding.

Note: You can also call with D3DGETDATA_FLUSH in dwGetDataFlags. This blocks the return of the GetData call until the occlusion query result has been returned. However, this approach slows your application noticeably.

```
ULONG IDirect3DQuery9::Release(void);
```

Be sure to free all objects that you create.

General Occlusion Query Usage

In general, most graphics cards buffer a few frames to allow some efficiency in the parallel work between the driver and the card. In order to use a single occlusion query, you must frame lock your application to buffer only one frame ahead (for example, flush). Forcing the driver into lockstep with the card will lower performance.

Because each occlusion query only represents one test, use multiple queries if you want to test multiple objects in a scene.

If your simple occlusion hull does not fully cover your complex object, the object will blink in and out of the complex object.

Steps to Use an Occlusion Query

The steps to use an occlusion query are as follows:

1. Render the scene
2. Turn off color and z writes
3. Begin the query
4. Render an “invisible” simple hull object that covers complex object
5. End the query
6. Test the results of the occlusion query
7. Turn color and z writes back on
8. If coarse hull is completely occluded, then the query is done
9. If coarse hull is not occluded, then render the complex object

Note: These steps show what needs to happen for a single occlusion query. If you are testing multiple objects, maintain a queue of queries to avoid putting the hardware and driver into lock-step.

LatentOcclusionQueryBank

Included with this demo is a class called LatentOcclusionQueryBank. This class wraps up a queue of occlusion queries to handle the standard frame queuing that occurs on most modern graphics hardware

API

HRESULT GetLatestResults (DWORD *count);

Returns S_OK if new results are returned since the last time GetLatestResults was called. The occlusion query fragment count is passed back in the count parameter.

HRESULT BeginNextQuery();

Issues the beginning of a new query. Returns S_FALSE if it has run out of queries.

HRESULT EndNextQuery();

Issues the end of a previously begun query. Returns S_FALSE if it has run out of queries or if it cannot find a previous query.

unsigned int GetNumActiveQueries();

Returns the number of queries that have not yet had results delivered. This number is updated when GetLatestResults is called.

void Free();

Frees all occlusion queries that have been allocated. Always call free when you are done.

API Mapping

LPDIRECT3DQUERY9	→	LatentOcclusionQueryBank
GetData()	→	GetLatestResults()
Issue(D3DISSUE_BEGIN)	→	BeginNextQuery()
Issue(D3DISSUE_BEGIN)	→	EndNextQuery()
Release()	→	Free()

Note: Replacing all calls to a single IDirect3DQuery9 object with the equivalent calls into LatentOcclusionQueryBank should work smoothly.



Notice

ALL NVIDIA DESIGN SPECIFICATIONS, REFERENCE BOARDS, FILES, DRAWINGS, DIAGNOSTICS, LISTS, AND OTHER DOCUMENTS (TOGETHER AND SEPARATELY, "MATERIALS") ARE BEING PROVIDED "AS IS." NVIDIA MAKES NO WARRANTIES, EXPRESSED, IMPLIED, STATUTORY, OR OTHERWISE WITH RESPECT TO THE MATERIALS, AND EXPRESSLY DISCLAIMS ALL IMPLIED WARRANTIES OF NONINFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE.

Information furnished is believed to be accurate and reliable. However, NVIDIA Corporation assumes no responsibility for the consequences of use of such information or for any infringement of patents or other rights of third parties that may result from its use. No license is granted by implication or otherwise under any patent or patent rights of NVIDIA Corporation. Specifications mentioned in this publication are subject to change without notice. This publication supersedes and replaces all information previously supplied. NVIDIA Corporation products are not authorized for use as critical components in life support devices or systems without express written approval of NVIDIA Corporation.

Trademarks

NVIDIA and the NVIDIA logo are trademarks or registered trademarks of NVIDIA Corporation in the United States and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

Copyright

© 2004 NVIDIA Corporation. All rights reserved



NVIDIA.

NVIDIA Corporation
2701 San Tomas Expressway
Santa Clara, CA 95050
www.nvidia.com